



RAJALAKSHMI ENGINEERING COLLEGE

Approved by AICTE | Affiliated to Anna University | Accredited by NAAC

Department of Computer Science and Engineering

CS23334 Fundamentals of Data Science Lab

III semester II Year (2023R)

Name of the Student : SARVESH R

Register Number : 240701478

202

5

nnjusjhrr

November 18,

```
[8]: from matplotlib import pyplot as plt
overs=[10,20,30,40,50] india=[64,106,135,191,254]
nz=[69,99,131,172,251]
plt.plot(overs,india,color='b',marker='o') plt.plot(overs,nz,color='black',marker='o')
plt.xlabel("scores") plt.ylabel("overs")
plt.title("ind vs nz champions trophy 2025")
plt.grid(True) plt.show()
```




202
5

[]:

ipaczigq7

November 18,

```
[15]: from matplotlib import pyplot as plt
matches=['VS bangladesh','VS Pakistan','VS New Zealand','VS Australia','VS Newzealand final']
scores=[231,244,249,267,254]
plt.plot(matches,scores,color='b',marker='o',linestyle='--')
plt.title("scores of india in champions trophy 2025")
plt.xlabel("MATCHES->")
plt.ylabel("scores")
plt.grid(True)
plt.show()
```




202
5

[]:

exp1c

18,

```
[3]: from matplotlib import pyplot as plt
ipl={'CSK':5,'MI':5,'KKR':3,'SRH':1,'RR':1,'GT':1,'DC':1,'RCB':1}
cups=list(ipl.values()) teams=list(ipl.keys())
colors=['yellow','blue','violet','orange','pink','black','green','red']
plt.bar(teams,cups,color=colors) plt.title("IPL CHAMPIONS",color='b')
plt.xlabel("teams") plt.ylabel("no of cups") plt.grid(True) plt.show()
```



November 2025

[]:

exp1d


18,

```
[8]: from matplotlib import pyplot as plt

student_data = {
    'A': (1, 50),
    'B': (2, 55),
    'C': (3, 65),
    'D': (4, 70),
    'E': (5, 75),
    'F': (6, 78),
    'G': (7, 85),
    'H': (8, 88),
    'I': (9, 93),
    'J': (10, 95)
}
students=list(student_data.keys())
hours=[value[0] for value in student_data.values()]
marks=[value[1] for value in student_data.values()]

plt.scatter(students,hours,color='g',marker='*',s=100)
plt.scatter(students,marks,color='b',marker='.',s=200)
plt.xlabel("names")
plt.ylabel("marks") plt.title("student
marks") plt.show()
```

November 2025



[]:


exp1e

18,

```
[11]: from matplotlib import pyplot as plt

ages = [18, 19, 20, 20, 21, 21, 22, 22, 22, 23,
        23, 24, 25, 25, 25, 26, 27, 28, 28, 29,
        30, 30, 30, 31, 32, 34, 35, 36, 38, 40]

plt.hist(ages,color='g',bins=10,edgecolor='black')
plt.xlabel("ages") plt.ylabel("count") plt.grid(True)
plt.show()
```




```
[ ]:
```

November 2025

exp1f

18,

```
[17]: from matplotlib import pyplot as plt  
  
brands = ['Samsung', 'Apple', 'Xiaomi', 'OnePlus', 'Oppo', 'Vivo', 'Others']  
market_share = [30, 25, 15, 10, 8, 7, 5]  
  
plt.pie(market_share, labels=brands, autopct='%.0f%%', colors=['red', 'green', 'blue', 'white', 'yellow', 'orange', 'violet']) plt.show()
```



November 2025

[]:

exp2

November 18, 2025

```
[4]: import os import pandas as pd import
      numpy as np from matplotlib import
      pyplot as plt os.chdir("D:\\pandas")
df=pd.read_csv("sales_data.csv")
df.head()
```

```
[4]:          Date    Product Sales Quantity Region
 0  01-01-2023  Product A    200     4   North
 1  02-01-2023  Product B    150     3   South
 2  03-01-2023  Product A    220     5   North
 3  04-01-2023  Product C    300     6   East
 4  05-01-2023  Product B    180     4   West
```

```
[7]: df.isnull().sum()
```

```
[7]: Date      0
      Product    0
      Sales      0
      Quantity    0
      Region      0
      dtype: int64
```

```
[13]: df['Sales'].fillna(df['Sales'].mean())
```

```
[13]: 0      200
 1      150
 2      220
```

```
3    300
4    180
5    210
6    320

7    160
8    230
9    310
10   190
11   240
12   330
13   170
14   250
15   340
Name: Sales, dtype: int64
```

```
[15]: df.dropna(subset=['Product','Quantity','Region'])
```

```
[15]:          Date      Product Sales Quantity Region
0  01-01-2023  Product A     200    4 North
1  02-01-2023  Product B     150    3 South
2  03-01-2023  Product A     220    5 North
3  04-01-2023  Product C     300    6 East
4  05-01-2023  Product B     180    4 West
5  06-01-2023  Product A     210    5 North
6  07-01-2023  Product C     320    7 East
7  08-01-2023  Product B     160    3 South
8  09-01-2023  Product A     230    6 North
9  10-01-2023  Product C     310    7 East
10 11-01-2023  Product B     190    4 West
11 12-01-2023  Product A     240    6 North
12 13-01-2023  Product C     330    8 East
```

```
13 14-01-2023 Product B      170  3 South
14 15-01-2023 Product A      250  7 North
15 16-01-2023 Product C      340  8     East
```

```
[16]: df.describe()
```

```
[16]: Sales  Quantity      count
16.000000 16.000000
    mean 237.500000  5.375000
    std   64.031242  1.746425
    min  150.000000  3.000000
    25% 187.500000  4.000000
    50% 225.000000  5.500000
    75% 302.500000  7.000000
    max 340.000000  8.000000
```

```
[18]: product_summary=df.groupby('Product').agg({
    'Sales':'sum',
    'Quantity':'sum'
}).reset_index()
```

```
product_summary
```


```
[18]:      Product Sales Quantity
0  Product A      1350  33
1  Product B      850   17
2  Product C      1600  36
```

```
[19]: plt.figure(figsize=(10,6))
```


```
[19]: <Figure size 1000x600 with 0 Axes>
```

```
<Figure size 1000x600 with 0 Axes>
```

```
[20]: plt.bar(product_summary['Product'], product_summary['Sales'])
plt.xlabel('Product')
plt.ylabel('Total Sales')
plt.title('Total Sales by Product')
plt.show()
```



```
[22]: sales_over_time = df.groupby('Date').agg({'Sales': 'sum'}).reset_index()
plt.figure(figsize=(10, 6))
plt.plot(sales_over_time['Date'],sales_over_time['Sales'])
plt.xlabel('Date') plt.ylabel('Total Sales') plt.title('SalesOver Time')
plt.show()
```



```
pivot_table = df.pivot_table(values='Sales', index='Region', columns='Product',
aggfunc='sum',fill_value=0)

pivot_table
```

```
[29]:
```

Product	Product A	Product B	Product C	Region
East	0	0	1600	
North	1350	0	0	

```
South      0      480      0
West       0      370      0
```

```
df['Date'] = pd.to_datetime(df['Date'], dayfirst=True)
correlation_matrix = df.select_dtypes(include=[float, int]).corr()
print(correlation_matrix)
```


[37]:

```
Sales Quantity
Sales    1.000000  0.944922
Quantity  0.944922  1.000000
```

[38]:

```
import seaborn as sns
plt.figure(figsize=(8, 6))

sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



[]:

exp3

November 18, 2025

```
[2]: import pandas as pd  
import numpy as np  
import os  
  
os.chdir("D:\\pandas")  
  
df=pd.read_csv("pre_process_datasample.csv")  
  
df
```

```
[2]:   Country    Age    Salary Purchased  
0     France  44.0  72000.0        No  
1      Spain  27.0  48000.0       Yes  
2   Germany  30.0  54000.0        No  
3      Spain  38.0  61000.0        No  
4   Germany  40.0      NaN       Yes  
5     France  35.0  58000.0       Yes  
6      Spain      NaN  52000.0        No  
7     France  48.0  79000.0       Yes  
8   Germany  50.0  83000.0        No  
9     France  37.0  67000.0       Yes
```

```
[3]: df["Country"].mode()
```

```
[3] : 0      France  
Name: Country, dtype: object
```

```
df.info()
```

```
[4] :
```

```
<class          'pandas.core.frame.DataFrame'>  
RangeIndex: 10 entries, 0 to 9
```

Data columns (total 4 columns):

#	Column	Non-Null Count	Dtype
0	Country	10 non-null	object
1	Age	9 non-null	float64
2	Salary	9 non-null	float64
3	Purchased	10 non-null	object

dtypes: float64(2), object(2)
memory usage: 452.0+ bytes

```
[5]: df["Country"].mode()[0]
```

```
[5]: 'France'
```

```
[6]: type(df["Country"].mode()[0])
```

```
[6]: str
```

```
[8]: df["Country"].fillna(df["Country"].mode()[0])
```

```
[8]: 0    France
      1    Spain
      2    Germany
      3    Spain
      4    Germany
      5    France
      6    Spain
      7    France
      8    Germany
      9    France
```

Name: Country, dtype: object

```
[10]: df["Age"].fillna(df["Age"].median())
```

```
[10]: 0    44.0
      1    27.0
      2    30.0
      3    38.0
      4    40.0
      5    35.0
```

```
6    38.0
7    48.0
8    50.0
9    37.0
Name: Age, dtype: float64
```

```
[48]: df["Salary"].fillna(round(df["Salary"].mean()))
```

```
[48]: 0      72000.0
1      48000.0
2      54000.0
3      61000.0
4      63778.0
5      58000.0
6      52000.0
78  79000.083000.0
9      67000.0
Name: Salary, dtype: float64
```

```
[49]: pd.get_dummies(df.Country)
```

```
[49]:   France Germany Spain
0    True    False  False
1   False    False  True
2   False    True  False
3   False    False  True
4   False    True  False
5    True    False  False
6   False    False  True
7    True    False  False
8   False    True  False
9    True    False  False
```

```
[50]: updated_dataset = pd.concat([pd.get_dummies(df.Country), df.iloc[:, [1, 2, 3]], axis=1])
```

```
updated_dataset
```

```
[52]: updated_dataset["Purchased"] = updated_dataset["Purchased"].replace(['No','Yes'], [0,1]).astype(int)
```

```
updated_dataset
```

```
[50]:   France Germany Spain    Age   Salary Purchased
      0   True    False  False  44.0  72000.0        No
      1  False    False   True  27.0  48000.0       Yes
      2  False     True  False  30.0  54000.0        No
      3  False    False   True  38.0  61000.0        No
      4  False     True  False  40.0  63778.0       Yes
      5   True    False  False  35.0  58000.0       Yes
      6  False    False   True    NaN  52000.0        No
      7   True    False  False  48.0  79000.0       Yes
      8  False     True  False  50.0  83000.0        No
      9   True    False  False  37.0  67000.0       Yes
```

```
[53]:
```

```
[53]:   France Germany Spain    Age   Salary Purchased
      0   True    False  False  44.0  72000.0    0
      1  False    False   True  27.0  48000.0    1
```

2 False True False 30.0 54000.0 0

3 False False True 38.0 61000.0 0

4 False True False 40.0 63778.0 1

5 True False False 35.0 58000.0 1

6 False False True NaN 52000.0 0

7 True False False 48.0 79000.0 1

8 False True False 50.0 83000.0 0

9 True False False 37.0 67000.0 1

[]:

exp4

November 18, 2025

```
[45]: import numpy as np  
import pandas as pd  
import os  
  
os.chdir("D:\\pandas")  
  
df=pd.read_csv("Hotel_Dataset.csv")  
df
```

```
[45]:   CustomerID Age_Group Rating(1-5)          Hotel FoodPreference Bill \\  
0           1    20-25            4      Ibis        veg 1300  
1           2    30-35            5  LemonTree     Non-Veg 2000  
2           3    25-30            6    RedFox        Veg 1322  
3           4    20-25           -1  LemonTree        Veg 1234  
4           5     35+             3      Ibis  Vegetarian  989  
5           6     35+             3     Ibys     Non-Veg 1909  
6           7     35+             4    RedFox  Vegetarian 1000  
7           8    20-25            7  LemonTree        Veg 2999  
8           9    25-30            2      Ibis     Non-Veg 3456  
9           9    25-30            2      Ibis     Non-Veg 3456
```

```
10      10    30-35      5   RedFox      non-Veg -6755
```

	NoOfPax	EstimatedSalary	Age_Group.1
0	2	40000	20-25
1	3	59000	30-35
2	2	30000	25-30
3	2	120000	20-25
4	2	45000	35+
5	2	122220	35+
6	-1	21122	35+
7	-10	345673	20-25
8	3	-99999	25-30
9	3	-99999	25-30
10	4	87777	30-35

```
[46]: df.duplicated()
```

```
[46]: 0    False
1    False
2    False
3    False
4    False
5    False
6    False
7    False
8    False
9    True
10   False dtype: bool
```

```
df.info()
```

```
[47]:
```

```
<class
'pandas.core.frame.DataFrame'> RangeIndex:
```

11 entries, 0 to 10 Data columns (total 9

columns):

#	Column	Non-Null Count	Dtype
0	CustomerIDAge_Group	1111	nonnon--nullnull int64object 1
2	Rating(1-5)	11	non-null int64
3	Hotel	11	non-null object
4	FoodPreference	11	non-null object
5	BillNoOfPax	1111	nonnon--nullnull int64int64 6
7	EstimatedSalary	11	non-null int64 8 Age_Group.1
11	non-null object		
	dtypes:	int64(5), object(4)	memory
	usage:	924.0+	bytes

```
df
```

[48] :

```
df.drop_duplicates(inplace=True)
```

[48]:

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill \
0	1	20-25	4	Ibis	veg	1300
1	2	30-35	5	LemonTree	Non-Veg	2000
2	3	25-30	6	RedFox	Veg	1322
3	4	20-25	-1	LemonTree	Veg	1234
4	5	35+ 3	Ibis	Vegetarian	989	
5	6	35+ 3	Ibys	Non-Veg	1909	
6	7	35+ 4	RedFox	Vegetarian	1000	
7	8	20-25	7	LemonTree	Veg	2999
8	9	25-30	2	Ibis	Non-Veg	3456
10	10	30-35	5	RedFox	non-Veg	-6755
		NoOfPax	EstimatedSalary	Age_Group.1		

```

0    2                  40000      20-25
1    3                  59000      30-35
2    2                  30000      25-30
3    2                  120000     20-25
4    2      5      2      45000      35+
6    -1                 122220     35+
7    -10                21122     35+
8    3                  345673     20-25
                    -99999     25-30
10       4                  87777     30-35

```

```
len(df)
```

[49] :

[49]: 10

```

index=np.array(list(range(0,len(df))))
df.set_index(index)

```

[50] :

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax
0	1	20-25	4	Ibis	veg	1300	2
1	2	30-35	5	LemonTree	Non-Veg	2000	3
2	3	25-30	6	RedFox	Veg	1322	3
3	4	20-25	-1	LemonTree	Veg	1234	3
4	5	35+	3	Ibis	Vegetarian	989	3
5	6	35+	3	Ibys	Non-Veg	1909	3
6	7	35+	4	RedFox	Vegetarian	1000	-1
7	8	20-25	7	LemonTree	Veg	2999	-10
8	9	25-30	2	Ibis	Non-Veg	3456	3
9	10	30-35	5	RedFox	non-Veg	-6755	4

EstimatedSalary Age_Group.1

0	40000	20-25
1	59000	30-35
2	30000	25-30
3	120000	20-25
4	45000	35+
5	122220	35+
6	21122	35+
7	345673	20-25
8	-99999	25-30
9	87777	30-35

```
[51]: df=df.drop(['Age_Group.1'],axis=1)
```

```
[51]:   CustomerID Age_Group Rating(1-5)          Hotel FoodPreference Bill \
0           1    20-25            4      Ibis        veg 1300
1           2    30-35            5  LemonTree    Non-Veg 2000
2           3    25-30            6     RedFox        Veg 1322
3           4    20-25           -1  LemonTree        Veg 1234
4           5    35+             3      Ibis  Vegetarian  989
5           6    35+             3     IbyS    Non-Veg 1909
6           7    35+             4     RedFox  Vegetarian 1000
7           8    20-25            7  LemonTree        Veg 2999
8           9    25-30            2      Ibis    Non-Veg 3456
```

10	10	30-35	5	RedFox	non-Veg	-6755
NoOfPax EstimatedSalary						
0	2	40000				
1	3	59000				
2	2	30000	3	2	120000	4
45000	5	2	122220			
6	-1	21122				
7	-10	345673				
8	3	-99999				
10	4	87777				

```
[52]: df.loc[df["CustomerID"] < 0, "CustomerID"] = np.nan
df.loc[df["Bill"] < 0, "Bill"] = np.nan
df.loc[df["EstimatedSalary"] < 0, "EstimatedSalary"] = np.nan
df.loc[df["Rating(1-5)"] < 0, "Rating(1-5)"] = np.nan
```

df

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill \
0	1.0	20-25	4.0	Ibis	veg	1300.0
1	2.0	30-35	5.0	LemonTree	Non-Veg	2000.0
2	3.0	25-30	6.0	RedFox	Veg	1322.0
3	4.0	20-25	Nan	LemonTree	Veg	1234.0
4	5.0	35+	3.0	Ibis	Vegetarian	989.0
5	6.0	35+	3.0	Ibys	Non-Veg	1909.0
6	7.0	35+	4.0	RedFox	Vegetarian	1000.0
7	8.0	20-25	7.0	LemonTree	Veg	2999.0
8	9.0	25-30	2.0	Ibis	Non-Veg	3456.0
10	10.0	30-35	5.0	RedFox	non-Veg	NaN
NoOfPax EstimatedSalary						
0	2	40000.0				
1	3	59000.0				

```
2      2 30000.0 3 2 120000.0 4 2 45000.0 5 2 122220.0 6 -1 21122.0
7      -10     345673.0
8      3 NaN
10     4     87777.0
```

```
[53]: df.Age_Group.unique()
```

```
[53]: array(['20-25', '30-35', '25-30', '35+'], dtype=object)
```

```
[54]: df["Hotel"].unique()
```

```
[54]: array(['Ibis', 'LemonTree', 'RedFox', 'Ibys'], dtype=object)
```

```
[55]: df.Hotel.replace(['Ibys'], 'Ibis')
```

```
[55]: 0      Ibis
1      LemonTree
2      RedFox
3      LemonTree
4      Ibis
5      Ibis
6      RedFox
7      LemonTree
8      Ibis
10     RedFox
Name: Hotel, dtype: object
```

```
[56]: df.FoodPreference.replace(['Vegetarian','veg'],'Veg')
```

```
[56]: 0      Veg
1      Non-
    Veg
2      Veg
3      Veg
4      Veg
5      Non-
    Veg
```

```
6      Veg
7      Veg
8  Non-
    Veg
10  non-
    Veg
```

Name: FoodPreference, dtype: object

```
[57]: df.FoodPreference.replace(['non-Veg'],'Non-Veg')
```

```
[57]: 0      veg    1
      Non-Veg
2      Veg
3      Veg  4      Vegetarian
5      Non-Veg
6      Vegetarian
7      Veg
8      Non-Veg
10     Non-Veg
```

Name: FoodPreference, dtype: object

```
[58]: df.EstimatedSalary.fillna(round(df.EstimatedSalary.mean()))
```

```
[58]: 0      40000.0
1      59000.0
2      30000.0
3      120000.0
4      45000.0
5      122220.0
6      21122.0
7      345673.0
8      96755.0
10     87777.0
```

Name: EstimatedSalary, dtype: float64

```
[59]: df.NoOfPax.fillna(round(df.NoOfPax.median()))
df['Rating(1-5)'].fillna(round(df['Rating(1-5)'].median()))
df.Bill.fillna(round(df.Bill.mean()))
df
```

```
[59]:    CustomerID  Age_Group Rating(1-5)      Hotel FoodPreference    Bill \
0           1.0    20-25          4.0       Ibis        veg 1300.0
1           2.0    30-35          5.0  LemonTree     Non-Veg 2000.0
2           3.0    25-30          6.0     RedFox        Veg 1322.0
3           4.0    20-25         NaN  LemonTree        Veg 1234.0
4           5.0    35+            3.0       Ibis  Vegetarian  989.0
5           6.0    35+            3.0      Ibys     Non-Veg 1909.0
6           7.0    35+            4.0     RedFox  Vegetarian 1000.0
7           8.0    20-25          7.0  LemonTree        Veg 2999.0
8           9.0    25-30          2.0       Ibis     Non-Veg 3456.0
10          10.0   30-35          5.0     RedFox    non-Veg     NaN
                                         NoOfPax EstimatedSalary
0           2        40000.0

```

```
1           3      59000.0
2           2      30000.0
3           2     120000.0
4           2     45000.0
5           2    122220.0
6           -1     21122.0
7          -10    345673.0
8           3      NaN
10          4      87777.0
```

```
[ ]:
```


exp6

November 18, 2025

```
[1]: import numpy as np
```

```
[2]: import pandas as pd
```

```
[7]: import os
```

```
os.chdir("D:\\pandas")
df=pd.read_csv('pre_process_datasample.csv')
```

```
[7]:   Country  Age  Salary Purchased
```

```
0  France  44.0  72000.0      No
1    Spain  27.0  48000.0     Yes
```

```
2 Germany  30.0  54000.0      No
```

```
3    Spain  38.0  61000.0      No
```

```
4 Germany  40.0        NaN     Yes
```

```
5    France  35.0  58000.0     Yes
```

```
6    Spain    NaN  52000.0      No
```

```
7    France  48.0  79000.0     Yes
```

```
8 Germany  50.0  83000.0      No
```

```
9    France  37.0  67000.0     Yes
```

```
[5]: df.head()
```

```
[5]:   Country  Age  Salary Purchased
```

```
0  France  44.0  72000.0      No
```

```
1    Spain  27.0  48000.0     Yes
```

```
2 Germany  30.0  54000.0      No
```

```
3    Spain  38.0  61000.0     No
```

```
4 Germany  40.0        NaN     Yes
```

```
[6]: df.Country.fillna(df.Country.mode()[0])
```

```
[6]: 0      France  1  
      Spain  
2      Germany  
3      Spain  
4      Germany  
5      France  
6      Spain  
7      GermanyFrance  
8  
9      France  
Name: Country, dtype: object
```

```
[8]: features=df.iloc[:,:-1].values
```

```
[ ]: label=df.iloc[:, -1].values
```

```
[10]: from sklearn.impute import SimpleImputer  
  
age=SimpleImputer(strategy="mean",missing_values=np.nan)  
  
Salary=SimpleImputer(strategy="mean",missing_values=np.nan)
```

```
[11]: age.fit(features[:,[1]])
```

```
[11] : SimpleImputer()
```

```
[12] : Salary.fit(features[:,[2]])
```

```
[12] : SimpleImputer()
```

```
[13] : SimpleImputer()
```

```
[13] : SimpleImputer()
```

```
features[:,[1]]=age.transform(features[:,[1]])  
features[:,[2]]=Salary.transform(features[:,[2]])  
features
```

```
[14] :
```

```
[14] : array(['France', 44.0, 72000.0],  
           ['Spain', 27.0, 48000.0],  
           ['Germany', 30.0, 54000.0],  
           ['Spain', 38.0, 61000.0],  
           ['Germany', 40.0, 63777.7777777778],  
           ['France', 35.0, 58000.0],  
           ['Spain', 38.77777777777778, 52000.0],  
           ['France', 48.0, 79000.0],  
           ['Germany', 50.0, 83000.0],  
           ['France', 37.0, 67000.0]], dtype=object)
```

```
[15] : from sklearn.preprocessing import OneHotEncoder oh =  
OneHotEncoder(sparse_output=False)  
Country=oh.fit_transform(features[:,0])
```

```
[15]: array([[1., 0., 0.],  
           [0., 0., 1.],  
           [0., 1., 0.],  
           [0., 0., 1.],  
           [0., 1., 0.],  
           [1., 0., 0.],  
           [0., 0., 1.],  
           [1., 0., 0.],  
           [0., 1., 0.],  
           [1., 0., 0.]])
```

```
[16] : final_set=np.concatenate((Country,features[:,1,2]),axis=1)  
final_set
```

```
[16]: array([[1.0, 0.0, 0.0, 44.0, 72000.0],  
           [0.0, 0.0, 1.0, 27.0, 48000.0],  
           [0.0, 1.0, 0.0, 30.0, 54000.0],  
           [0.0, 0.0, 1.0, 38.0, 61000.0],  
           [0.0, 1.0, 0.0, 40.0, 63777.7777777778],
```

```
[1.0, 0.0, 0.0, 35.0, 58000.0],  
[0.0, 0.0, 1.0, 38.77777777777778, 52000.0],  
[1.0, 0.0, 0.0, 48.0, 79000.0],  
[0.0, 1.0, 0.0, 50.0, 83000.0],  
[1.0, 0.0, 0.0, 37.0, 67000.0]], dtype=object)
```

```
[17]: from sklearn.preprocessing import StandardScaler  
sc=StandardScaler()  
sc.fit(final_set)  
feat_standard_scaler=sc.transform(final_set)  
feat_standard_scaler
```

```
[17]: array([[ 1.22474487e+00, -6.54653671e-01, -6.54653671e-01,  
    7.58874362e-01, 7.49473254e-01],  
   [-8.16496581e-01, -6.54653671e-01, 1.52752523e+00,  
    -1.71150388e+00, -1.43817841e+00],  
   [-8.16496581e-01, 1.52752523e+00, -6.54653671e-01,  
    -1.27555478e+00, -8.91265492e-01],  
   [-8.16496581e-01, -6.54653671e-01, 1.52752523e+00,  
    -1.13023841e-01, -2.53200424e-01],  
   [-8.16496581e-01, 1.52752523e+00, -6.54653671e-01,  
    1.77608893e-01, 6.63219199e-16],  
   [ 1.22474487e+00, -6.54653671e-01, -6.54653671e-01,  
    -5.48972942e-01, -5.26656882e-01],  
   [-8.16496581e-01, -6.54653671e-01, 1.52752523e+00,  
    0.0000000e+00, -1.07356980e+00],  
   [ 1.22474487e+00, -6.54653671e-01, -6.54653671e-01, 1.34013983e+00,  
    1.38753832e+00],  
   [-8.16496581e-01, 1.52752523e+00, -6.54653671e-01,  
    1.63077256e+00, 1.75214693e+00],  
   [ 1.22474487e+00, -6.54653671e-01, -6.54653671e-01,  
    -2.58340208e-01, 2.93712492e-01]])
```

```
[18]: from sklearn.preprocessing import MinMaxScaler  
mms=MinMaxScaler(feature_range=(0,1))  
mms.fit(final_set)  
feat_minmax_scaler=mms.transform(final_set)  
feat_minmax_scaler
```

```
[18]: array([[1.        , 0.        , 0.        , 0.73913043, 0.68571429],  
           [0.        , 0.        , 1.        , 0.        , 0.        ],  
           [0.        , 1.        , 0.        , 0.13043478, 0.17142857],  
           [0.        , 0.        , 1.        , 0.47826087, 0.37142857],  
           [0.        , 1.        , 0.        , 0.56521739, 0.45079365],  
           [1.        , 0.        , 0.        , 0.34782609, 0.28571429],  
           [0.        , 0.        , 1.        , 0.51207729, 0.11428571],  
           [1.        , 0.        , 0.        , 0.91304348, 0.88571429],  
           [0.        , 1.        , 0.        , 1.        , 1.        ],  
           [1.        , 0.        , 0.        , 0.43478261, 0.54285714]])
```

```
[ ]:
```


exp7

November 18, 2025


```
[1]: import seaborn as sns  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
%matplotlib inline  
tips=sns.load_dataset('tips')
```

```
[2]: tips.head()
```

```
[2]:   total_bill  tip    sex smoker  day    time  size  
0      16.99  1.01  Female     No  Sun  Dinner     2  
1      10.34  1.66    Male     No  Sun  Dinner     3  
2      21.01  3.50    Male     No  Sun  Dinner     3  
3      23.68  3.31    Male     No  Sun  Dinner     2  
4      24.59  3.61  Female     No  Sun  Dinner     4
```

```
[3]: sns.displot(tips.total_bill,kde=True)
```


```
[3] : <seaborn.axisgrid.FacetGrid at 0x201fdd6a660>
```



[4] :

```
sns.displot(tips.total_bill,kde=False)
```


[4] : <seaborn.axisgrid.FacetGrid at 0x201fdf4b250>



[5] :

```
sns.jointplot(x=tips.tip,y=tips.total_bill)
```


[5] : <seaborn.axisgrid.JointGrid at 0x201fdde9fd0>



[6] :

```
sns.jointplot(x=tips.tip,y=tips.total_bill,kind="reg")
```


[6] : <seaborn.axisgrid.JointGrid at 0x201ff15c550>



[7] :

```
sns.jointplot(x=tips.tip,y=tips.total_bill,kind="hex")
```


[7] : <seaborn.axisgrid.JointGrid at 0x201ff56da90>



[8] :

```
sns.pairplot(tips)
```

[8] : <seaborn.axisgrid.PairGrid at 0x201fddeaf90>



[9] :

```
tips.time.value_counts()
```

[9]: time Dinner


176

Lunch 68 Name:
count, dtype: int64

[10]:

```
sns.pairplot(tips,hue='time')
```


[10]: <seaborn.axisgrid.PairGrid at 0x201843e9310>



[11]:


```
sns.pairplot(tips,hue='day')
```

[11]: <seaborn.axisgrid.PairGrid at 0x20184cbd950>



```
[12]: sns.heatmap(tips.corr(numeric_only=True), annot=True)
```


```
[12]: <Axes: >
```



[13]:

```
sns.boxplot(tips.total_bill)
```


[13]: <Axes: ylabel='total_bill'>



[14]:

```
sns.boxplot(tips.tip)
```


[14]: <Axes: ylabel='tip'>



[15]:

```
sns.countplot(tips.day)
```


[15]: <Axes: xlabel='count', ylabel='day'>



[16]:

```
sns.countplot(tips.sex)
```


[16]: <Axes: xlabel='count', ylabel='sex'>



[17]:

```
tips.sex.value_counts().plot(kind='bar')
```


[17]: <Axes: ylabel='count'>



[18]:

```
tips.sex.value_counts().plot(kind='bar')
```


[18]: <Axes: xlabel='sex'>



[19]:

```
sns.countplot(tips[tips.time=='Dinner']['day'])
```

[19]: <Axes: xlabel='count', ylabel='day'>



[]:

```
In [ ]: import numpy as np
import pandas as pd
df=pd.read_csv('Salary_data.csv')
df
```

```
In [19]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   YearsExperience  30 non-null    float64
 1   Salary            30 non-null    int64  
dtypes: float64(1), int64(1)
memory usage: 612.0 bytes
```

```
In [3]: df.dropna(inplace=True)
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   YearsExperience  30 non-null    float64
 1   Salary            30 non-null    int64  
dtypes: float64(1), int64(1)
memory usage: 612.0 bytes
```

```
In [5]: df.describe()
```

```
Out[5]:
```

	YearsExperience	Salary
count	30.000000	30.000000
mean	5.313333	76003.000000
std	2.837888	27414.429785
min	1.100000	37731.000000
25%	3.200000	56720.750000
50%	4.700000	65237.000000
75%	7.700000	100544.750000
max	10.500000	122391.000000

```
In [6]: features=df.iloc[:,[0]].values
label=df.iloc[:,[1]].values
```

```
In [7]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(features,label,test_size=0.2,random_st
```

```
In [20]: from sklearn.linear_model import LinearRegression
model=LinearRegression()
model.fit(x_train,y_train)
```

```
In [21]: model.score(x_train,y_train)
```

```
Out[21]: 0.9603182547438908
```

```
In [23]: model.score(x_test,y_test)
```

```
Out[23]: 0.9184170849214232
```

```
In [24]: model.coef_
```

```
Out[24]: array([[9281.30847068]])
```

```
In [25]: model.intercept_
```

```
Out[25]: array([27166.73682891])
```

```
In [26]: import pickle  
pickle.dump(model,open('SalaryPred.model','wb'))
```

```
In [27]: model=pickle.load(open('SalaryPred.model','rb'))
```

```
In [28]: yr_of_exp=float(input("Enter Years of Experience: "))  
yr_of_exp_NP=np.array([[yr_of_exp]])  
Salary=model.predict(yr_of_exp_NP)
```

```
Enter Years of Experience: 44
```

```
In [ ]:
```

```
In [29]: print("Estimated Salary for {} years of experience is {}:".format(yr_of_exp,Salary))
```

```
Estimated Salary for 44.0 years of experience is [[435544.30953887]]:
```

```
In [ ]:
```

```
In [1]: import numpy as np  
import pandas as pd  
df=pd.read_csv('Social_Network_Ads.csv')  
df
```

```
Out[1]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
..
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

400 rows × 5 columns

```
In [2]: df.head()
```

```
Out[2]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
In [4]: features=df.iloc[:,[2,3]].values  
label=df.iloc[:,4].values  
features
```

```
In [51]: label
```

```
In [6]: from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LogisticRegression
```

```
In [7]: for i in range(1,401):
    x_train,x_test,y_train,y_test=train_test_split(features,label,test_size=0.
    model=LogisticRegression()
    model.fit(x_train,y_train)
    train_score=model.score(x_train,y_train)
    test_score=model.score(x_test,y_test)
    if test_score>train_score:
        print("Test {} Train{} Random State {}".format(test_score,train_score,i))
```

```
Test 0.6875 Train0.63125 Random State 3
Test 0.7375 Train0.61875 Random State 4
Test 0.6625 Train0.6375 Random State 5
Test 0.65 Train0.640625 Random State 6
Test 0.675 Train0.634375 Random State 7
Test 0.675 Train0.634375 Random State 8
Test 0.65 Train0.640625 Random State 10
Test 0.6625 Train0.6375 Random State 11
Test 0.7125 Train0.625 Random State 13
Test 0.675 Train0.634375 Random State 16
Test 0.7 Train0.628125 Random State 17
Test 0.7 Train0.628125 Random State 21
Test 0.65 Train0.640625 Random State 24
Test 0.6625 Train0.6375 Random State 25
Test 0.75 Train0.615625 Random State 26
Test 0.675 Train0.634375 Random State 27
Test 0.7 Train0.628125 Random State 28
Test 0.6875 Train0.63125 Random State 29
Test 0.6875 Train0.63125 Random State 31
```

```
In [8]: x_train,x_test,y_train,y_test=train_test_split(features,label,test_size=0.2,  
finalModel=LogisticRegression()  
finalModel.fit(x_train,y_train)
```

Out[8]: LogisticRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [9]: print(finalModel.score(x_train,y_train))
        print(finalModel.score(x_test,y_test))
```

0.834375
0.9125

```
In [10]: from sklearn.metrics import classification_report  
print(classification_report(label.finalModel.predict(features)))
```

	precision	recall	f1-score	support
0	0.85	0.93	0.89	257
1	0.84	0.71	0.77	143

```
In [1]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
%matplotlib inline
```

```
In [2]: df=pd.read_csv('Mall_Customers.csv')
```

```
In [3]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 200 entries, 0 to 199  
Data columns (total 5 columns):  
 #   Column           Non-Null Count  Dtype     
 ---  --  
 0   CustomerID      200 non-null    int64    
 1   Gender          200 non-null    object    
 2   Age             200 non-null    int64    
 3   Annual Income (k$) 200 non-null    int64    
 4   Spending Score (1-100) 200 non-null    int64    
 dtypes: int64(4), object(1)  
 memory usage: 7.9+ KB
```

```
In [4]: df.head()
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
In [5]: sns.pairplot(df)
```

```
Out[5]: <seaborn.axisgrid.PairGrid at 0x170e8e47850>
```



```
In [6]: features=df.iloc[:,[3,4]].values
```

```
In [7]: from sklearn.cluster import KMeans
model=KMeans(n_clusters=5)
model.fit(features)
KMeans(n_clusters=5)

C:\Users\Ayyadurai\AppData\Local\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of 'n_init' will change from 10 to 'auto' in 1.4. Set the value of 'n_init' explicitly to suppress the warning
    warnings.warn(
C:\Users\Ayyadurai\AppData\Local\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
    warnings.warn(
```

```
Out[7]: KMeans(n_clusters=5)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [8]: Final=df.iloc[:,[3,4]]
Final['label']=model.predict(features)
Final.head()
```

```
C:\Users\Ayyadurai\AppData\Local\Temp\ipykernel_8116\470183701.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
Final['label']=model.predict(features)
```

```
Out[8]:   Annual Income (k$)  Spending Score (1-100)  label
```

	Annual Income (k\$)	Spending Score (1-100)	label
0	15	39	4
1	15	81	2
2	16	6	4
3	16	77	2
4	17	40	4

```
In [9]: sns.set_style("whitegrid")
sns.FacetGrid(Final,hue="label",height=8) \
.map(plt.scatter,"Annual Income (k$)", "Spending Score (1-100)") \
.add_legend();
plt.show()
```



```
In [10]: features_el=df.iloc[:,[2,3,4]].values
from sklearn.cluster import KMeans
wcss=[]
for i in range(1,10):
    model=KMeans(n_clusters=i)
    model.fit(features_el)
    wcss.append(model.inertia_)
plt.plot(range(1,10),wcss)
```